

CLOUD COMPUTING PROJECT


WATCHES

WEBSERVICES

PART II

Part II - Objectives

- Deploy Watch Service in **Google Cloud**
- Leverage PaaS
 - GCloud Kubernetes
 - GCloud SQL
 - Docker Images repository
 - NoSQL DB (AWS S3)
- Automate the process
 - Code → Deployment



Microservice - Image info v1

- New Image (proxy) microservice
 - One endpoint `/image/v1/get/{sku}`
 - Will transfer an image from a NoSQL repo (AWS S3) to the user
 - Fixed mime type (image/png)
 - No image processing (ATM ;-)
 - Set proper caching headers and ETag
 - Lifetime 1h
 - OpenAPI description in git: `image_openapi_v1.yaml`

Watches Images

- The images are stored in a AWS S3 bucket (Ireland)
 - Publicly accessible
 - You can access them from GKE
 - Image format is transparent PNG (max 1024x1024)
 - <https://s3-eu-west-1.amazonaws.com/cloudcomputing-2018/project1/images/<SKU>.png>
 - <https://s3-eu-west-1.amazonaws.com/cloudcomputing-2018/project1/images/CAC1111.BA0850.png>
 - <https://s3-eu-west-1.amazonaws.com/cloudcomputing-2018/project1/images/CV201AP.FC6429.png>
 - Not all images will be available ! (~ 50% of SKUs have images)
 - Check for 200 or 404 HTTP codes

DockerHub

- Docker/Kubernetes/AWS can use any Docker images repositories
 - DockerHub is the default
 - If you publish your images (publicly) on DockerHub, you will not have to manage complex access to a particular repo (and no authentication to manage)
 - Create a free account
- Microservices images to publish
 - Infos v1
 - Images v1
- Publishing updated images should be part of your build process

Kubernetes - GKE

- Create a Google Cloud Platform account
 - <https://cloud.google.com/compute/>
 - 300\$ / 1 year free credits
 - Create a Kubernetes cluster in Europe West (close to AWS S3 for images)
 - Link your account to kubectl
 - <https://cloud.google.com/kubernetes-engine/docs/quickstart>
 - `$ gcloud container clusters get-credentials <cluster> --zone europe-west1-b --project <project>`
 - Then, use kubectl as you will use it locally with minikube !
- Kubernetes also exists on AWS but is expensive (~200\$ / month with minimal config) and much complicated to deploy
 - Kubernetes is a technology from Google and is better integrated in GCE



Home



Billing



Kubernetes Engine



SQL



Storage



PRODUCTS



Marketplace



Billing



APIs & Services



Support



IAM & admin



Getting started



Security



COMPUTE

DASHBOARD

ACTIVITY

CUSTOMIZE



Project info



Project name

kube-2019

Project ID

kube-2019-256711

Project number

938662733845



Go to project settings



Resources



Compute Engine

1 instance



Storage

1 bucket



SQL

1 instance



Trace



No trace data from the past 7 days



Compute Engine



CPU (%)



instance/cpu/utilization: 0.204




Go to Compute Engine



SQL



Storage used (bytes)



Google Cloud Platform status



All services normal



Go to Cloud status dashboard



Billing



Estimated charges

USD \$0.00

For the billing period Oct 1 – 23, 2019



View detailed charges



Error Reporting



No sign of any errors. Have you set up Error Reporting?



Learn how to set up Error Reporting



News



Use GKE usage metering to combat over-provisioning

[←](#) Create a Kubernetes cluster

Clone an existing cluster

Select one of your existing clusters to populate fields



Standard cluster

Continuous integration, web serving, backends. Best choice for further customization or if you are not sure what to choose.



Your first cluster

Experimenting with Kubernetes Engine, deploying your first application. Affordable choice to get started. [→](#)



CPU intensive applications

Web crawling or anything else that requires more CPU.



Memory intensive applications

Databases, analytics, things like Hadoop, Spark, ETL or anything else that requires more memory.



GPU Accelerated Computing

Machine learning, video transcoding, scientific computations or anything else that is compute-intensive and can utilize GPUs.



Highly available

Most demanding availability requirements. Both the master and the nodes are replicated across multiple zones.

Name [?]Location type [?]☒ Zonal☐ RegionalZone [?]

Master version

ⁱ Try the new Release Channels feature instead of managing the master version directly.

[Use Release Channels](#)

Node pools

Node pools are separate instance groups running Kubernetes in a cluster. You may add node pools in different zones for higher availability, or add node pools of different type machines. To add a node pool, click Edit. [Learn more](#)

pool-1

Number of nodes

Machine configuration [?]

Machine family

☒ General-purpose☐ Memory-optimized

Machine types for common workloads, optimized for cost and flexibility

Generation

First

Powered by Skylake CPU platform or one of its predecessors

Machine type



vCPU

1 shared core

Memory

1.7 GB

Cluster version	1.14.6-gke.13 (latest)
Machine type	g1-small
Autoscaling	Disabled
Stackdriver Logging & Monitoring	Disabled
Boot disk size	30GB

You will be billed for the 1 node (VM instance) in your cluster. [Compute Engine pricing](#)

- Clusters
- Workloads
- Services & Ingress
- Applications
- Configuration
- Storage

A Kubernetes cluster is a managed group of VM instances for running containerized applications. [Learn more](#)

<input type="checkbox"/>	Name ^	Location	Cluster size	Total cores	Total memory	Notifications	Labels	
<input type="checkbox"/>	test-2019	europe-west1-b	1	1 vCPU	1.70 GB			Connect

```
leo@t480s: ~  
gcloud container clusters get-credentials test-2019 --zone europe-west1-b --project kube-2019-256711  
Fetching cluster endpoint and auth data.  
kubeconfig entry generated for test-2019.  
~$ kubectl cluster-info  
Kubernetes master is running at https://35.195.108.21  
GLBCDefaultBackend is running at https://35.195.108.21/api/v1/namespaces/kube-system/services/default-http-backend:http/proxy  
Heapster is running at https://35.195.108.21/api/v1/namespaces/kube-system/services/heapster/proxy  
KubeDNS is running at https://35.195.108.21/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy  
Metrics-server is running at https://35.195.108.21/api/v1/namespaces/kube-system/services/https:metrics-server:/proxy  
To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.  
~$
```

- Clusters
- Workloads
- Services & Ingress
- Applications
- Configuration
- Storage

test-2019

Details Storage Nodes

Nodes

Filter nodes

Name ^	Status	CPU requested	CPU allocatable	Memory requested	Memory allocatable
gke-test-2019-pool-1-bb16ecb7-zf6b	Ready	629 mCPU	940 mCPU	634.8 MB	1.22 GB

SQL Database

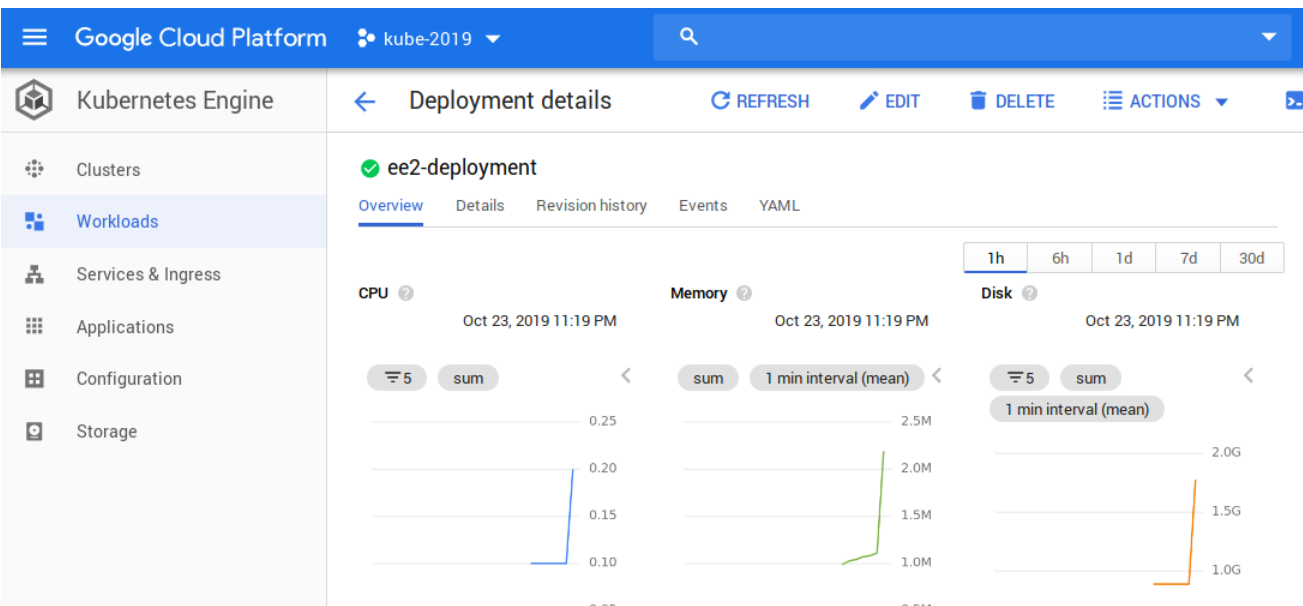
- The SQL database will be used by info services v1
 - Deploy the same database as in part I
- The product to use is Google Cloud SQL
 - Managed MySQL
 - See doc below to access it from Kubernetes pods
 - <https://cloud.google.com/sql/docs/mysql/connect-kubernetes-engine>

Kubernetes Architecture

- **Objects to create** (at the end all have to be described in a single yaml file)
 - **Deployments**
 - One deployment for each docker image (→ 2)
 - 2 replicas (pods)
 - **Services**
 - One service for each image deployment (→ 2)
 - Service will load balance between the pods of the corresponding deployment
 - **Ingress**
 - Will route to the right service depending on the path

Kubernetes - Ingress controller

- Use ingress controller to route to your microservices
 - `/info/v1/*` → Service info
 - `/image/v1/*` → Service image
- Use **HTTP** (HTTPS not mandatory)
- <https://cloud.google.com/kubernetes-engine/docs/concepts/ingress>
- <https://cloud.google.com/kubernetes-engine/docs/tutorials/http-balancer>
- Note: you should wait a few (10-15) minutes before the ingress controller is really deployed



Google Cloud Platform kube-2019

Kubernetes Engine

Services & Ingress

REFRESH CREATE INGRESS DELETE

Kubernetes services Brokered services BETA Ingresses

Services are sets of Pods with a network endpoint that can be used for discovery and load balancing. Ingresses are collections of rules for routing external HTTP(S) traffic to Services.

Is system object : False Filter resources

Name	Status	Type	Endpoints	Pods	Namespace	Cluster
ee2-ingress	Ok	Ingress	35.201.69.5/	0 / 0	default	test-2019
ee2-service	Ok	Load balancer	35.195.207.231:80	2 / 2	default	test-2019

Deliverables

- `/info/`
 - Move here all files from project part I
- `/image/`
 - Write the new service (follow project part I guidelines):
 - Dockerfile, server.py, requirements.txt, run.sh
- `/README`
 - DNS:port to test your API !
 - Additional indications, specially if something is not working as expected
- `/build.sh`
 - Rebuild the 2 docker service images
- `/deploy.sh`
 - Publish/update the 2 docker images to DockerHub
 - 2 instances for each services (replicas: 2)
 - Kubernetes rolling upgrade
- `/all.yaml`
 - Deployments + Services + Ingress
 - Multiple object descriptions can be grouped in one file, use '\n---\n' as separator

Delay - Grading

- Delay: 4 weeks
 - Deadline: 2019-11-20T23:59:59+02:00
- Grading
 - Participation
 - 1 point
 - Endpoint with working APIs for Info & Image
 - 3 points
 - Scripts to automate build & rolling upgrade
 - 2 points

Documentation

- Minikube
 - <https://github.com/kubernetes/minikube>
 - <https://kubernetes.io/docs/tasks/tools/install-minikube/>
 - <https://kubernetes.io/docs/setup/learning-environment/minikube/>
- Kubernetes Cheat Sheet
 - <https://linuxacademy.com/blog/containers/kubernetes-cheat-sheet/>